Learning Multi-Step Predictive State Representations

Lucas Langer (♣), Borja Balle (♦), Doina Precup (♣) McGill University (♣), Lancaster University (♦)

July 12, 2016

Predictive State Representations (PSR)

Motivation

- Make predictions in partially observable environments
- 2 Learn a representation of hidden states

Advantages over HMMs

- Statistically consistent and well understood learning efficiency
- Smaller representations of state [Boots et al., 2011; Hsu et al., 2009; Bailly et al., 2010]

Why Multi-Step PSRs?

- Leverage structure of underlying environment
- 2 Directly apply transitions which occur at longer time-scales

PSR: The single observation case

- PSR defined by: $\langle \alpha_0, \{A_\sigma\}, \alpha_\infty \rangle$, where α_0 is an initial weighting on states (1xn) A_σ is a transition matrix (nxn) α_∞ is a normalizer (nx1) n: number of latent states
- ② PSRs compute probabilities of observations $f(\sigma^k) = \alpha_0 \cdot A_{\sigma}^k \cdot \alpha_{\infty}$

[Littman et al., 2001; Singh et al., 2004; Rosencrantz et al., 2004]

Spectral Learning of PSRs

- Represent Data as a Hankel Matrix
- Singular Value Decomposition
- Pick Model Size
- Linear Algebra
- **5** Result: $\langle \alpha_0, \{A_\sigma\}, \alpha_\infty \rangle$

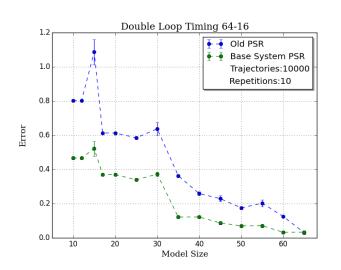
[Boots et al., 2011]

Adding multi-step operators: The Base System

• Learn $\{A_{\sigma^2}, A_{\sigma^4}, A_{\sigma^8}, ... A_{\sigma^{2N}}\}$ as extra transition operators Operators learned separately so $A_{\sigma} \cdot A_{\sigma} \neq A_{\sigma^2}$

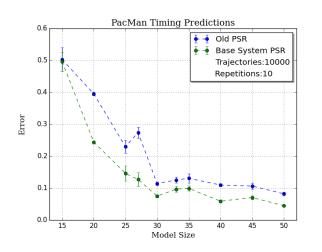
Why might this help? Reduce build up of error Faster computations

Base System Performance for Loops



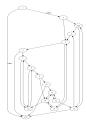


Pacman-like Labyrinth





(a) Pacman



(b) Graph

In general, which operators to learn?

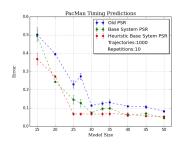
- Observations: $\{a^{30}b^{15}, a^{60}, b^{15}\}$ Desired operators: $A_{a^{30}b^{15}}, A_{a^{60}}, A_{a^{30}}, A_{b^{15}}, A_a, A_b$
- Sub-string properties: long, frequent, diverse Structured environments should be easier
- Trade-off between number of operators and learning time Iterative greedy heuristic works well
 Could also try an entropy based approach

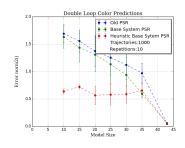
How should we execute queries

Minimize number of matrices
 Represent transitions as compactly as possible

- ② Query string: "abcacb", Operators = $\{A_{ab}, A_{bca}, A_{cb}, A_a, A_b\}$ Desired partition: "a—bca—cb" Computation: $f(abcacb) = \alpha_0 \cdot A_a \cdot A_{bca} \cdot A_{cb} \cdot \alpha_{\infty}$
- Solution: dynamic programming
 State update for online applications

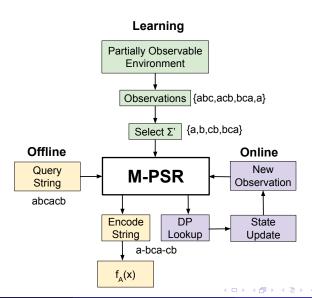
Performance of Heuristics







The Big Picture



Questions?