

Learning Multi-Step Predictive State Representations

Lucas Langer, Borja Balle, Doina Precup

July 10, 2016

Predictive State Representations (PSR)

Goals:

- 1 Make predictions in partially observable environments
- 2 Learn a representation of hidden states

Motivation:

- 1 Globally optimal predictions (unlike HMMs)
- 2 Learn smaller representations

Why Multi-Step PSRs?

- 1 Leverage structure at different time-scales
- 2 Learn state transitions for interesting observation sequences

PSR: The single observation case

- 1 PSR defined by: $\langle \alpha_0, \{A_\sigma\}, \alpha_\infty \rangle$

where

α_0 is an initial weighting on states $1 \times n$

A_σ is a transition matrix $n \times n$

α_∞ is a normalizer $n \times 1$

- 2 PSRs compute probabilities of observations

$$f(\sigma^k) = \alpha_0 \cdot A_\sigma^k \cdot \alpha_\infty$$

Spectral Learning of PSRs

Step 1: Represent Data as a Hankel Matrix

Step 2: Singular Value Decomposition

Step 3: Pick Model Size

Step 4: Linear Algebra

Result: $\langle \alpha_0, \{A_\sigma\}, \alpha_\infty \rangle$

Adding multi-step operators: The Base System

- Idea: Learn $\{A_\sigma, A_{\sigma^2}, A_{\sigma^4}, A_{\sigma^8}, \dots A_{\sigma^{2N}}\}$ as extra transition operators

Note: operators learned separately

$$f(\sigma^{11}) = \alpha_0 \cdot A_{\sigma^8} \cdot A_{\sigma^2} \cdot A_{\sigma^1} \cdot \alpha_\infty$$

- **Why might this help?**


Reduce error build up

Capture structure


Faster computations

Base System Performance for Loops

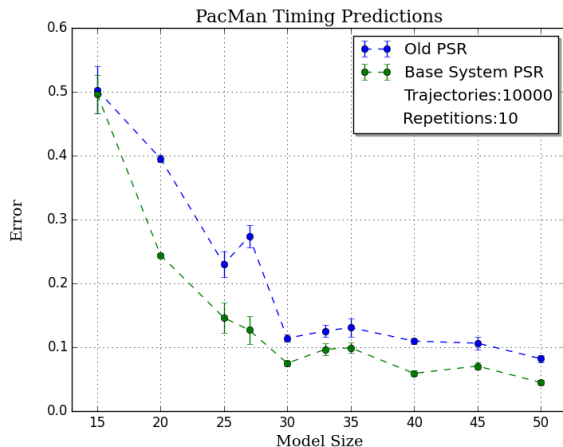
64-16 Loop Lengths




47-27 Loop Lengths



Pacman-like Labyrinth





(a) Pacman



(b) Graph

Wall Color Predictions

We paint the first loop green and the second loop blue



In general, which operators to learn?

- Observations: $\{a^{30}:10, a^{60}:5, b^{18}:15\}$
Desired Operators: $A_{a^{30}}, A_{b^{18}}, A_a, A_b$
- Substring properties: **long, frequent, diverse**
Structured environments should be easier
- Iterative greedy heuristic works well
Could also try an entropy based approach

How should we execute queries

- **Minimize number of matrices**

Represent transitions as compactly as possible

- Query string: "abcacb", Operators = $\{A_{ab}, A_{bca}, A_{cb}, A_a, A_b\}$





Desired partition: "a—bca—cb"

Computation: $f(abcacb) = \alpha_0 \cdot A_a \cdot A_{bca} \cdot A_{cb} \cdot \alpha_\infty$


- Solution: dynamic programming

State update for online applications

Performance of Heuristics



The Big Picture



Questions?